

IE 7275 Project

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import SimpleImputer, IterativeImputer
from sklearn.preprocessing import MinMaxScaler
#from category_encoders import BinaryEncoder
from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, recall_score, plot_roc_curve
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier
from xgboost.sklearn import XGBClassifier
from sklearn.model_selection import GridSearchCV
from sklearn import svm
import seaborn as sns
```

```
In [2]: df_mhealth = pd.read_csv('survey.csv')
df_mhealth.head()
```

Out[2]:

	Timestamp	Age	Gender	Country	state	self_employed	family_history	treatment	work_interfere	no_employees
0	2014-08-27 11:29:31	37	Female	United States	IL	NaN	No	Yes	Often	6-2
1	2014-08-27 11:29:37	44	M	United States	IN	NaN	No	No	Rarely	More than 100
2	2014-08-27 11:29:44	32	Male	Canada	NaN	NaN	No	No	Rarely	6-2
3	2014-08-27 11:29:46	31	Male	United Kingdom	NaN	NaN	Yes	Yes	Often	26-100
4	2014-08-27 11:30:22	31	Male	United States	TX	NaN	No	No	Never	100-500

5 rows × 27 columns

```
In [3]: df_mhealth.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1259 entries, 0 to 1258
Data columns (total 27 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Timestamp            1259 non-null  object
1   Age                  1259 non-null  int64
2   Gender               1259 non-null  object
3   Country              1259 non-null  object
```

4	state	744	non-null	object
5	self_employed	1241	non-null	object
6	family_history	1259	non-null	object
7	treatment	1259	non-null	object
8	work_interfere	995	non-null	object
9	no_employees	1259	non-null	object
10	remote_work	1259	non-null	object
11	tech_company	1259	non-null	object
12	benefits	1259	non-null	object
13	care_options	1259	non-null	object
14	wellness_program	1259	non-null	object
15	seek_help	1259	non-null	object
16	anonymity	1259	non-null	object
17	leave	1259	non-null	object
18	mental_health_consequence	1259	non-null	object
19	phys_health_consequence	1259	non-null	object
20	coworkers	1259	non-null	object
21	supervisor	1259	non-null	object
22	mental_health_interview	1259	non-null	object
23	phys_health_interview	1259	non-null	object
24	mental_vs_physical	1259	non-null	object
25	obs_consequence	1259	non-null	object
26	comments	164	non-null	object

dtypes: int64(1), object(26)
memory usage: 265.7+ KB

Data Cleaning

In [4]: `df_mhealth.isna().sum()`

Out[4]:

Timestamp	0
Age	0
Gender	0
Country	0
state	515
self_employed	18
family_history	0
treatment	0
work_interfere	264
no_employees	0
remote_work	0
tech_company	0
benefits	0
care_options	0
wellness_program	0
seek_help	0
anonymity	0
leave	0
mental_health_consequence	0
phys_health_consequence	0
coworkers	0
supervisor	0
mental_health_interview	0
phys_health_interview	0
mental_vs_physical	0
obs_consequence	0
comments	1095

dtype: int64

In [5]: `df_mhealth.drop(columns=['Timestamp', 'state', 'comments'], inplace = True)`

In [6]: `df_mhealth['Age'].unique()`

```
Out[6]: array([[ 37, 44, 32, 31, 33,
 35, 39, 42, 23, 29,
 36, 27, 46, 41, 34,
 30, 40, 38, 50, 24,
 18, 28, 26, 22, 19,
 25, 45, 21, -29, 43,
 56, 60, 54, 329, 55,
99999999999, 48, 20, 57, 58,
 47, 62, 51, 65, 49,
-1726, 5, 53, 61, 8,
 11, -1, 72], dtype=int64)
```

```
In [7]: df_mhealth['Age'].replace([df_mhealth['Age'][df_mhealth['Age'] < 18], np.nan, inplace = True)
df_mhealth['Age'].replace([df_mhealth['Age'][df_mhealth['Age'] > 90]], np.nan, inplace = True)
```

```
In [8]: df_mhealth['Gender'].unique()
```

```
Out[8]: array(['Female', 'M', 'Male', 'male', 'female', 'm', 'Male-ish', 'maile',
'Trans-female', 'Cis Female', 'F', 'something kinda male?',
'Cis Male', 'Woman', 'f', 'Mal', 'Male (CIS)', 'queer/she/they',
'non-binary', 'Femake', 'woman', 'Make', 'Nah', 'All', 'Enby',
'fluid', 'Genderqueer', 'Female ', 'Androgynous', 'Agender',
'cis-female/femme', 'Guy (-ish) ^_^', 'male leaning androgynous',
'Male ', 'Man', 'Trans woman', 'msle', 'Neuter', 'Female (trans)',
'queer', 'Female (cis)', 'Mail', 'cis male', 'A little about you',
'Malr', 'p', 'femal', 'Cis Man',
'ostensibly male, unsure what that really means'], dtype=object)
```

```
In [9]: df_mhealth['Gender'].replace(['Male ', 'male', 'm', 'M', 'Male', 'Man', 'Cis Male', 'cis m',
'Cis Man', 'msle', 'Malr', 'Mal', 'maile', 'Make', ], 'Male', inplace = True)

df_mhealth['Gender'].replace(['Female ', 'female', 'F', 'f', 'Woman', 'woman', 'Female', 'f',
'cis-female/femme', 'Femake', ], 'Female', inplace = True)

df_mhealth["Gender"].replace(['Female (trans)', 'queer/she/they', 'non-binary', 'fluid', 'c',
'male leaning androgynous', 'Agender', 'A little about you', 'Nah', 'All', 'ostensibly mal',
'Genderqueer', 'Enby', 'p', 'Neuter', 'something kinda male?', 'Guy (-ish) ^_^', 'Trans wor
```

```
In [10]: male_country = df_mhealth[df_mhealth['Gender'] == 'Male'][['Country', 'Gender']]
female_country = df_mhealth[df_mhealth['Gender'] == 'Female'][['Country', 'Gender']]
male_country = male_country.value_counts()
female_country = female_country.value_counts()

male_country = pd.DataFrame(male_country).reset_index().rename(columns={0:'count'}).head(10)
female_country = pd.DataFrame(female_country).reset_index().rename(columns={0:'count'}).head(10)
male_country['count'] = male_country['count'] * -1
```

```
In [11]: import plotly.graph_objs as go
import plotly
fig = plotly.tools.make_subplots(
    rows=2,
    shared_yaxes=True,
    horizontal_spacing=0)

fig.append_trace(go.Bar(
    y = df_mhealth["Gender"].value_counts(),
    x = ["Male", "Female", "Others"],
    textfont = dict(size = 10, color = '#6aa87b'),
    textposition = 'outside',
    name = 'Count of Employees by Gender',
```

```

marker_color='#528B8B',
orientation = 'v'),
row=1, col=1)

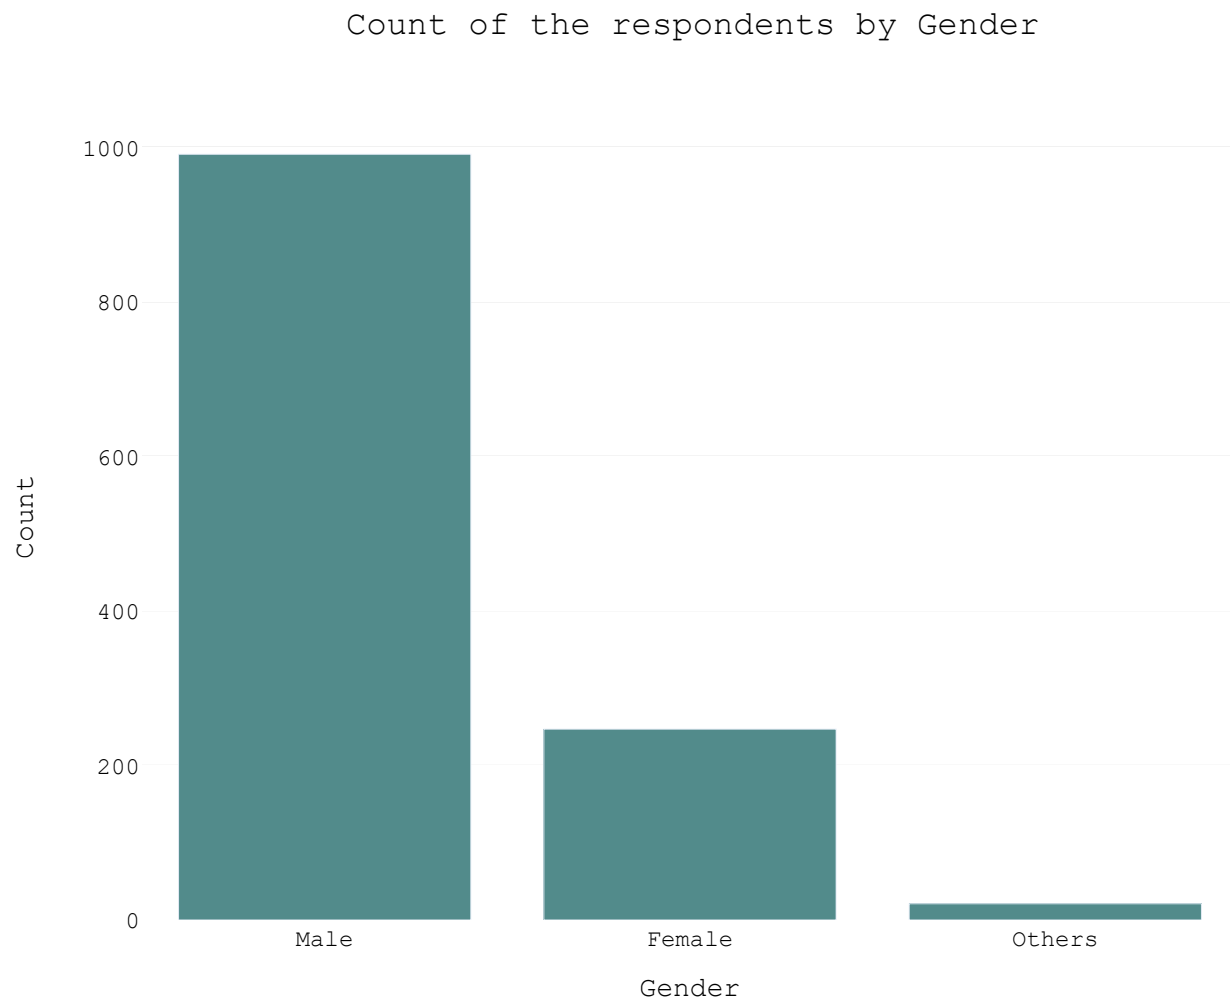
fig.update_layout(
    font_family = 'monospace',
    title = dict(text = 'Count of the respondents by Gender', x = 0.5,
margin = dict(t=80, b=0, l=70, r=40),
    hovermode = "y unified",
    plot_bgcolor = '#edf2c7',
    paper_bgcolor = '#edf2c7',
    xaxis_title = "Gender",
    yaxis_title = "Count",
    font = dict(color='black'),
    legend = dict(orientation="h",
        yanchor="bottom", y=1,
        xanchor="center", x=0.5),
    hoverlabel = dict(bgcolor="#edf2c7", font_size=13,
        font_family="Monospace"))

fig.show()

```

c:\users\amogha shettar\appdata\local\programs\python\python39\lib\site-packages\plotly\tools.py:461: DeprecationWarning:

plotly.tools.make_subplots is deprecated, please use plotly.subplots.make_subplots instead



In [12]:

```

import plotly.graph_objs as go
import plotly
fig = plotly.tools.make_subplots(rows=1, cols=2,

```

```

        specs=[[{}], {}]],
        shared_yaxes=True,
        horizontal_spacing=0)

fig.append_trace(go.Bar(
    y = male_country.Country,
    x = male_country['count'],
    text = male_country['count'],
    textfont = dict(size = 10, color = '#6aa87b'),
    textposition = 'outside',
    name = 'Male responses',
    marker_color='#EE6A50',
    orientation = 'h'),
    row=1, col=1)

fig.append_trace(go.Bar(
    y = male_country.Country,
    x = female_country['count'],
    text = female_country['count'],
    textfont = dict(size = 10, color = '#913f3f'),
    textposition = 'outside',
    name = 'Female responses',
    marker_color='#EEAD0E',
    orientation = 'h'),
    row=1, col=2)

fig.update_xaxes(
    tickfont = dict(size=15),
    tickmode = 'array',
    ticklen = 6,
    showline = False,
    showgrid = False,
    ticks = 'outside')

fig.update_yaxes(showgrid=False,
    categoryorder='total ascending',
    ticksuffix=' ',
    showline=False)

fig.update_layout(
    font_family = 'monospace',
    title = dict(text = 'Gender of the respondents across Countries',
    margin = dict(t=80, b=0, l=70, r=40),
    hovermode = "y unified",
    plot_bgcolor = '#edf2c7',
    paper_bgcolor = '#edf2c7',
    font = dict(color='black'),
    legend = dict(orientation="h",
        yanchor="bottom", y=1,
        xanchor="center", x=0.5),
    hoverlabel = dict(bgcolor="#edf2c7", font_size=13,
        font_family="Monospace"))


fig.show()

```

c:\users\amogha shettar\appdata\local\programs\python\python39\lib\site-packages\plotly\tols.py:461: DeprecationWarning:

plotly.tools.make_subplots is deprecated, please use plotly.subplots.make_subplots instead

Gender of the respondents across Countries



In [13]:

```
print('Maximum Age: ', df_mhealth.Age.max())
df_mhealth = df_mhealth[(df_mhealth.Age > 11) & (df_mhealth.Age <= 100)]
print(df_mhealth.Age.describe().astype(int))
```

```
Maximum Age: 72.0
count      1251
mean         32
std           7
min          18
25%          27
50%          31
75%          36
max          72
Name: Age, dtype: int32
```

In [14]:

```
fig1 = plotly.subplots.make_subplots(
    shared_yaxes=True,
    horizontal_spacing=0)

fig1.append_trace(go.Bar(
    y = df_mhealth["Age"].value_counts(),
    x = [29.0, 32.0, 26.0, 27.0, 33.0, 28.0, 31.0, 34.0, 30.0, 25.0, 35.0, 23.0, 24.0, 36.0],
    textfont = dict(size = 10, color = '#6aa87b'),
    textposition = 'outside',
    marker_color='#9ACD32',
    orientation = 'v'),
    row=1, col=1)

fig1.update_layout(
    font_family = 'monospace',
    title = dict(text = 'Age distribution of the respondents', x = 0, y = 0, align = 'left'),
    margin = dict(t=80, b=0, l=70, r=40),
    hovermode = "y unified",
    xaxis_title = "Age",
```


```

yaxis_title = "Count",
plot_bgcolor = '#edf2c7',
paper_bgcolor = '#edf2c7',
font
= dict(color='black'),
legend
= dict(orientation="h",
      yanchor="bottom", y=1,
      xanchor="center", x=0.5),
hoverlabel
= dict(bgcolor="#edf2c7", font_size=13,
      font_family="Monospace"))

```

```
fig1.show()
```

Age distribution of the respondents



```

In [15]: seek = df_mhealth[df_mhealth.treatment == 'Yes'].drop(['treatment', 'Country', 'Age'], axis=1)
          dont = df_mhealth[df_mhealth.treatment == 'No'].drop(['treatment', 'Country', 'Age'], axis=1)

```

```

In [16]: buttons = []
          i = 0
          vis = [False] * 21

          for col in seek.columns:
              vis[i] = True
              buttons.append({'label' : col,
                             'method' : 'update',
                             'args' : [{'visible' : vis},
                                       {'title' : col}] })

              i+=1
              vis = [False] * 21

          fig = plotly.tools.make_subplots(rows=1, cols=2,

```

```

        specs=[[{'type':'domain'}, {'type':'domain'}]])

for col in dont.columns:
    fig.add_trace(go.Pie(
        values = dont[col].value_counts(),
        labels = dont[col].value_counts().index,
        title = dict(text = 'No Treatment: <br>Distribution<br>of {}'.format(col),
            font = dict(size=18, family = 'monospace'),
        ),
        hoverinfo='label+percent',,1,1)

for col in seek.columns:
    fig.add_trace(go.Pie(
        values = seek[col].value_counts(),
        labels = seek[col].value_counts().index,
        title = dict(text = 'Seek Treatment: <br>Distribution<br>of {}'.format(col),
            font = dict(size=18, family = 'monospace'),
        ),
        hoverinfo='label+percent',,1,2)

fig.update_traces(hoverinfo='label+percent',
    textinfo='label+percent',
    textfont_size=12,
    opacity = 0.8,
    showlegend = False,
    marker = dict(colors = sns.color_palette('Blues').as_hex(),
        line=dict(color='#000000', width=1)))


fig.update_traces(row=1, col=2, hoverinfo='label+percent',
    textinfo='label+percent',
    textfont_size=12,
    opacity = 0.8,
    showlegend = False,
    marker = dict(colors = sns.color_palette('Oranges').as_hex(),
        line=dict(color='#000000', width=1)))

fig.update_layout(margin=dict(t=0, b=0, l=0, r=0),
    font_family = 'monospace',
    hovermode = "y unified",
    plot_bgcolor = '#edf2c7',
    paper_bgcolor = '#edf2c7',
    updatemenus = [dict(
        type = 'dropdown',
        x = 0.60,
        y = 0.95,
        showactive = True,
        active = 0,
        buttons = buttons)],
    annotations=[
        dict(text = "<b>Choose<br>Column<b> : ",
            font = dict(size = 14),
            showarrow=False,
            x = 0.5, y = 1.03, yref = "paper", align = "left")])

for i in range(1,42):
    fig.data[i].visible = False
fig.data[21].visible = True

fig.show()

```

In [17]:

```
adf = df_mhealth[(df_mhealth.anonymity == 'Yes') | (df_mhealth.anonymity == 'No')]
df_cross = pd.crosstab(adf.treatment,adf.anonymity)
# initiate data list for figure
data = []
#use for loop on every zoo name to create bar data
for x in df_cross.columns:
    data.append(go.Bar(name=str(x), x=df_cross.index, y=df_cross[x]))

figure = go.Figure(data)
figure.update_layout(barmode = 'group')

figure.update_xaxes(
    tickfont = dict(size=15),
    tickmode = 'array',
    ticklen = 6,
    showline = False,
    showgrid = False,
    ticks = 'outside')

figure.update_yaxes(showgrid=False,
    categoryorder='total ascending',
    ticksuffix=' ',
    showline=False)

figure.update_layout(
    font_family = 'monospace',
    title = dict(text = 'Is anonymity protected if a person decides',
    margin = dict(t=80, b=0, l=70, r=40),
    hovermode = "y unified",
    plot_bgcolor = '#edf2c7',
    paper_bgcolor = '#edf2c7',
    xaxis_title = "Anonymity",
    yaxis_title = "Count",
    legend_title = "Treatment",
```

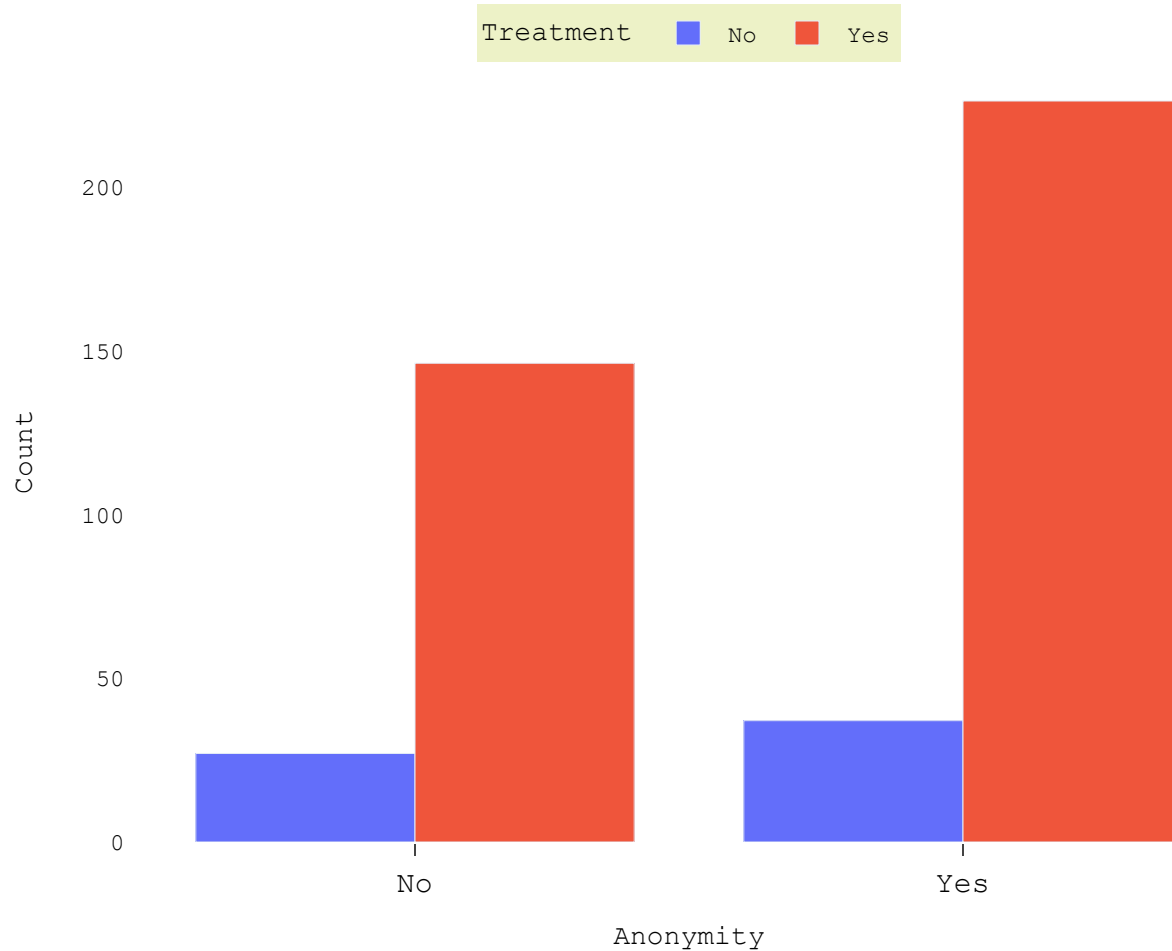
```

font          = dict(color='black'),
legend        = dict(orientation='h',
                      yanchor='bottom', y=1,
                      xanchor='center', x=0.5),
hoverlabel    = dict(bgcolor='#edf2c7', font_size=13,
                      font_family='Monospace'))

```

```
figure.show()
```

onymity protected if a person decides to take advantage of mental



In [18]:

```

fig = plotly.tools.make_subplots(rows=2, cols=1,
                                  shared_yaxes=True,
                                  horizontal_spacing=0)

fig.append_trace(go.Bar(
    y = df_mhealth.mental_health_interview.value_counts(),
    x = ["No", "Maybe", "Yes"],
    textfont = dict(size = 10, color = '#6aa87b'),
    textposition = 'outside',
    name = 'Mental Health',
    marker_color='#FF7F24',
    orientation = 'v'),
    row=1, col=1)

fig.append_trace(go.Bar(
    y = df_mhealth.phys_health_interview.value_counts(),
    x = ["Maybe", "No", "Yes"],
    textfont = dict(size = 10, color = '#913f3f'),
    textposition = 'outside',
    name = 'Physical Health',
    marker_color='#8B8B83',
    orientation = 'v'),
    row=2, col=1)

```

```

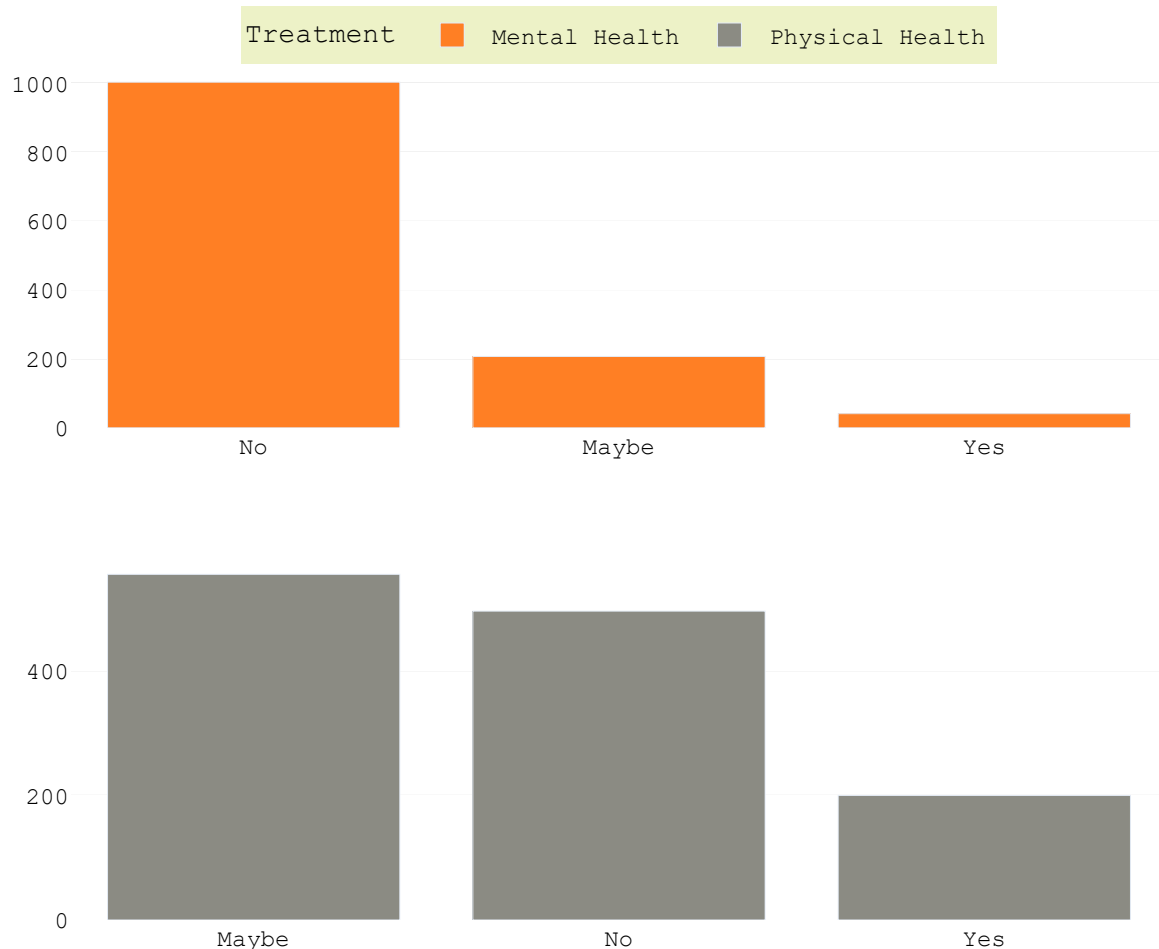
row=2, col=1)

fig.update_layout(
    font_family = 'monospace',
    title = dict(text = 'Attitude of the respondants towards Mental
margin = dict(t=80, b=0, l=70, r=40),
    hovermode = "y unified",
    plot_bgcolor = '#edf2c7',
    paper_bgcolor = '#edf2c7',
    #xaxis_title = "Response",
    #yaxis_title = "Count",
    legend_title = "Treatment",
    font = dict(color='black'),
    legend = dict(orientation="h",
                    yanchor="bottom", y=1,
                    xanchor="center", x=0.5),
    hoverlabel = dict(bgcolor="#edf2c7", font_size=13,
                      font_family="Monospace"))

fig.show()

```

Attitude of the respondants towards Mental and Physical health



```
In [19]: df_mhealth.drop(columns=['Country'], inplace = True)
```

Preprocessing

```
In [20]: mode_onehot_pipe = Pipeline([
    ('encoder', SimpleImputer(strategy = 'most_frequent')),
    ('one hot encoder', OneHotEncoder(handle_unknown = 'ignore'))])

```

```
transformer = ColumnTransformer([
    ('one_hot', OneHotEncoder(handle_unknown = 'ignore'), ['Gender', 'family_history', 'no',
                                                         'remote_work', 'tech_company',
                                                         'wellness_program', 'anonymity',
                                                         'mental_health_interview', 'phy',
                                                         'mental_vs_physical', 'obs_consequence', 'mental_health_consequence',
                                                         'seek_help']),
    ('mode_onehot_pipe', mode_onehot_pipe, ['self_employed', 'work_interfere']),
    ('iterative', IterativeImputer(max_iter = 10, random_state = 0), ['Age'])])
```

```
In [21]: df_mhealth['treatment'] = np.where(df_mhealth['treatment'] == 'Yes', 1, 0)
```

```
In [22]: X = df_mhealth.drop(df_mhealth[['treatment']], axis = 1)
y = df_mhealth['treatment']
```

```
In [23]: X_train, X_test, y_train, y_test = train_test_split(X, y, stratify = y, test_size = 0.3, r
```

```
In [24]: def model_evaluation(model, metric):
    model_cv = cross_val_score(model, X_train, y_train, cv = StratifiedKFold(n_splits = 5))
    return model_cv
```

```
In [25]: log_model = LogisticRegression()
```

```
In [26]: log_pipe = Pipeline([('transformer', transformer), ('log_model', log_model)])

log_pipe_cv = model_evaluation(log_pipe, 'recall')

for model in [log_pipe]:
    model.fit(X_train, y_train)

log_score_cv = [log_pipe_cv.round(2)]
log_score_recall_score = [recall_score(y_test, log_pipe.predict(X_test))
                           ]
method_name = ['Logistic Regression']
log_cv = pd.DataFrame({
    'method': method_name,
    'cv score': log_score_cv,
    'recall score': log_score_recall_score
})
log_cv
```

```
Out[26]:
```

	method	cv score	recall score
0	Logistic Regression	[0.69, 0.77, 0.74, 0.76, 0.72]	0.742105

```
In [27]: decision_model = DecisionTreeClassifier(random_state = 1111)
```

```
In [28]: decision_pipe = Pipeline([('transformer', transformer), ('decision_model', decision_model)])

decision_pipe_cv = model_evaluation(decision_pipe, 'recall')

for model in [decision_pipe]:
    model.fit(X_train, y_train)
```

```

decision_score_cv = [decision_pipe_cv.round(2)]
decision_score_recall_score = [recall_score(y_test, decision_pipe.predict(X_test))
                                ]
method_name = ['Decision Tree Classifier']
decision_cv = pd.DataFrame({
    'method': method_name,
    'cv score': decision_score_cv,
    'recall score': decision_score_recall_score
})
decision_cv

```

Out[28]:

	method	cv score	recall score
0	Decision Tree Classifier	[0.63, 0.65, 0.67, 0.61, 0.66]	0.678947

In [29]:

```

svm_model = svm.SVC()

```

In [30]:

```

svm_pipe = Pipeline([('transformer', transformer), ('svm_model', svm_model)])

svm_pipe_cv = model_evaluation(svm_pipe, 'recall')

for model in [svm_pipe]:
    model.fit(X_train, y_train)

svm_score_cv = [svm_pipe_cv.round(2)]
svm_score_recall_score = [recall_score(y_test, svm_pipe.predict(X_test))
                           ]
method_name = ['Support Vector Machine']
svm_cv = pd.DataFrame({
    'method': method_name,
    'cv score': svm_score_cv,
    'recall score': svm_score_recall_score
})
svm_cv

```

Out[30]:

	method	cv score	recall score
0	Support Vector Machine	[0.65, 0.67, 0.7, 0.66, 0.65]	0.563158

In [31]:

```

random_model = RandomForestClassifier(random_state = 1111)

```

In [32]:

```

random_pipe = Pipeline([('transformer', transformer), ('random_model', random_model)])

random_pipe_cv = model_evaluation(random_pipe, 'recall')

for model in [random_pipe]:
    model.fit(X_train, y_train)

random_score_cv = [random_pipe_cv.round(2)]
random_score_recall_score = [recall_score(y_test, random_pipe.predict(X_test))
                              ]
method_name = ['Random Forest Classifier']
random_cv = pd.DataFrame({
    'method': method_name,
    'cv score': random_score_cv,
    'recall score': random_score_recall_score
})
random_cv

```

```
Out[32]:
```

	method	cv score	recall score
0	Random Forest Classifier	[0.74, 0.76, 0.76, 0.7, 0.73]	0.715789

```
In [33]: #Encoding data
from sklearn import preprocessing
from sklearn.preprocessing import binarize, LabelEncoder, MinMaxScaler

labelDict = {}
for feature in df_mhealth:
    le = preprocessing.LabelEncoder()
    le.fit(df_mhealth[feature])
    le_name_mapping = dict(zip(le.classes_, le.transform(le.classes_)))
    df_mhealth[feature] = le.transform(df_mhealth[feature])
    labelKey = 'label_' + feature
    labelValue = [*le_name_mapping]
    labelDict[labelKey] = labelValue
```

```
In [34]: scaler = MinMaxScaler()
df_mhealth['Age'] = scaler.fit_transform(df_mhealth[['Age']])
```

```
In [35]: feature_all = ['Age', 'Gender', 'family_history', 'benefits', 'care_options', 'anonymity',
feature_cols = ['Age', 'Gender', 'family_history', 'benefits', 'care_options', 'anonymity',
                'no_employees', 'remote_work', 'tech_company', 'wellness_program', 'anonymity',
                'mental_health_interview', 'phys_health_interview', 'mental_vs_physical',
                'mental_health_consequence', 'phys_health_consequence', 'leave', 'self_employment']

X2 = df_mhealth[feature_cols]
X1 = df_mhealth[feature_all]
y1 = df_mhealth.treatment

X_train1, X_test1, y_train1, y_test1 = train_test_split(X1, y1, test_size=0.30, random_state=42)
X_train2, X_test2, y_train2, y_test2 = train_test_split(X2, y1, test_size=0.30, random_state=42)

# Create dictionaries for final graph//

# Use: methodDict['Stacking'] = accuracy_score
methodDict = {}
rmseDict = ()
```

```
In [36]: def evalClassModel(model, y_test2, y_pred_class2, plot=False):
    #Classification accuracy: percentage of correct predictions
    # calculate accuracy
    print('Accuracy:', metrics.accuracy_score(y_test2, y_pred_class2))

    confusion = metrics.confusion_matrix(y_test2, y_pred_class2)
    #[row, column]
    TP = confusion[1, 1]
    TN = confusion[0, 0]
    FP = confusion[0, 1]
    FN = confusion[1, 0]

    sns.heatmap(confusion,annot=True,fmt="d")
    plt.title('Confusion Matrix')
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.show()

    accuracy = metrics.accuracy_score(y_test2, y_pred_class2)
```

```

print('Error:', 1 - metrics.accuracy_score(y_test2, y_pred_class2))

false_positive_rate = FP / float(TN + FP)

#print('False Positive Rate:', false_positive_rate)

#Precision: When a positive value is predicted, how often is the prediction correct?
print('Precision:', metrics.precision_score(y_test2, y_pred_class2))

# IMPORTANT: first argument is true values, second argument is predicted probabilities
print('AUC Score:', metrics.roc_auc_score(y_test2, y_pred_class2))

# calculate cross-validated AUC
print('Cross-validated AUC:', cross_val_score(model, X2, y1, cv=10, scoring='roc_auc'))

model.predict_proba(X_test2)[0:10, 1]

y_pred_prob2 = model.predict_proba(X_test2)[: , 1]

y_pred_prob2 = y_pred_prob2.reshape(-1,1)
y_pred_class2 = binarize(y_pred_prob2)[0]

roc_auc = metrics.roc_auc_score(y_test2, y_pred_prob2)

predict_mine2 = np.where(y_pred_prob2 > 0.50, 1, 0)
confusion = metrics.confusion_matrix(y_test2, predict_mine2)

return accuracy

```

In [37]:

```

import sklearn.metrics as metrics

def logisticRegression2():
    logreg2 = LogisticRegression(C=60)
    logreg2.fit(X_train2, y_train2)

    # make class predictions for the testing set
    y_pred_class2 = logreg2.predict(X_test2)

    print('Logistic Regression before feature selection -')

    accuracy_score = evalClassModel(logreg2, y_test2, y_pred_class2)


    #Data for final graph
    methodDict['Log. Regres.'] = accuracy_score * 100

```

In [38]:

```
logisticRegression2()
```

Logistic Regression before feature selection -
Accuracy: 0.7340425531914894



Error: 0.26595744680851063
Precision: 0.7790697674418605
AUC Score: 0.736281179138322
Cross-validated AUC: 0.7635268722772028

In [39]:

```
import sklearn.metrics as metrics

def SVM_C2():
    svm_c2 = svm.SVC(probability = True)
    svm_c2.fit(X_train2, y_train2)

    # make class predictions for the testing set
    y_pred_class2 = svm_c2.predict(X_test2)

    print('Support Vector Machine before feature selection -')


    accuracy_score = evalClassModel(svm_c2, y_test2, y_pred_class2)

    #Data for final graph
    #methodDict['Log. Regres.'] = accuracy_score * 100
```

In [40]:

```
SVM_C2()
```

Support Vector Machine before feature selection -
Accuracy: 0.7952127659574468



Error: 0.20478723404255317
Precision: 0.771689497716895

AUC Score: 0.7922335600907029
Cross-validated AUC: 0.8614626090168132

In [41]:

```
def randomForest2():
    # Calculating the best parameters
    forest2 = RandomForestClassifier(n_estimators = 100)

    featuresSize = feature_cols.__len__()
    param_dist = {"max_depth": [3, None],
                  # "max_features": randint(1, featuresSize),
                  # "min_samples_split": randint(2, 9),
                  # "min_samples_leaf": randint(1, 9),
                  "criterion": ["gini", "entropy"]}
    #tuningRandomizedSearchCV(forest2, param_dist)

    forest2 = RandomForestClassifier(max_depth = None, min_samples_leaf=8, min_samples_sp
    my_forest2 = forest2.fit(X_train2, y_train2)


    y_pred_class2 = my_forest2.predict(X_test2)

    print('Random Forest before feature selection-')
    accuracy_score = evalClassModel(my_forest2, y_test2, y_pred_class2, True)
```

In [42]:

```
randomForest2()
```

Random Forest before feature selection-
Accuracy: 0.7845744680851063



Error: 0.21542553191489366
Precision: 0.7601809954751131
AUC Score: 0.7813492063492063
Cross-validated AUC: 0.8726061287594538

In [43]:

```
from sklearn.ensemble import ExtraTreesClassifier

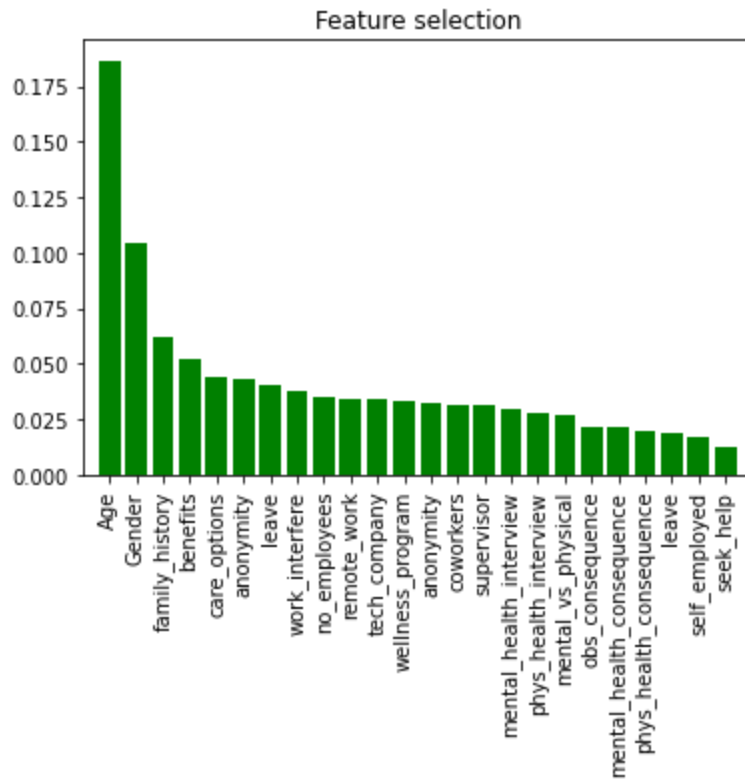
forest = ExtraTreesClassifier(n_estimators=10, random_state=1111)

forest.fit(X2, y1)
importances = forest.feature_importances_
std = np.std([tree.feature_importances_ for tree in forest.estimators_], axis=0)
indices = np.argsort(importances)[-1:]

labels = []
for f in range(X2.shape[1]):
```

```
labels.append(feature_cols[f])
```

```
#plt.figure(figsize=(12,24))
plt.title("Feature selection")
plt.bar(range(X2.shape[1]), importances[indices], color="g", align="center")
plt.xticks(range(X2.shape[1]), labels, rotation='vertical')
plt.xlim([-1, X2.shape[1]])
plt.show()
```



In [44]:

```
def evalClassModel(model, y_test1, y_pred_class, plot=False):
    #Classification accuracy: percentage of correct predictions
    # calculate accuracy
    print('Accuracy:', metrics.accuracy_score(y_test1, y_pred_class))

    confusion = metrics.confusion_matrix(y_test1, y_pred_class)
    #[row, column]
    TP = confusion[1, 1]
    TN = confusion[0, 0]
    FP = confusion[0, 1]
    FN = confusion[1, 0]

    sns.heatmap(confusion,annot=True,fmt="d")
    plt.title('Confusion Matrix')
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.show()

    accuracy = metrics.accuracy_score(y_test1, y_pred_class)

    print('Error:', 1 - metrics.accuracy_score(y_test1, y_pred_class))

    false_positive_rate = FP / float(TN + FP)

    #print('False Positive Rate:', false_positive_rate)

    #Precision: When a positive value is predicted, how often is the prediction correct?
    print('Precision:', metrics.precision_score(y_test1, y_pred_class))
```

```

# IMPORTANT: first argument is true values, second argument is predicted probabilities
print('AUC Score:', metrics.roc_auc_score(y_test1, y_pred_class))

# calculate cross-validated AUC
print('Cross-validated AUC:', cross_val_score(model, X1, y1, cv=10, scoring='roc_auc'))

#model.predict_proba(X_test1)[0:10, 1]

y_pred_prob = model.predict_proba(X_test1)[: , 1]

y_pred_prob = y_pred_prob.reshape(-1,1)

roc_auc = metrics.roc_auc_score(y_test1, y_pred_prob)

predict_mine = np.where(y_pred_prob > 0.50, 1, 0)
confusion = metrics.confusion_matrix(y_test1, predict_mine)

return accuracy

```

In [45]:

```

import sklearn.metrics as metrics

def logisticRegression():
    logreg = LogisticRegression(C=60)
    logreg.fit(X_train1, y_train1)

    # make class predictions for the testing set
    y_pred_class = logreg.predict(X_test1)

    print('Logistic Regression after feature selection -')

    accuracy_score = evalClassModel(logreg, y_test1, y_pred_class)


    #Data for final graph
    #methodDict['Log. Regres.']= accuracy_score * 100

```

In [46]:

```
logisticRegression()
```

Logistic Regression after feature selection -
Accuracy: 0.7180851063829787



Error: 0.28191489361702127
Precision: 0.7419354838709677

AUC Score: 0.7187074829931973
Cross-validated AUC: 0.7547748891994661

In [47]:

```
import sklearn.metrics as metrics

def SVM_C():
    svm_c = svm.SVC(probability = True)
    svm_c.fit(X_train1, y_train1)

    # make class predictions for the testing set
    y_pred_class = svm_c.predict(X_test1)

    print('Support Vector Machine after feature selection -')


    accuracy_score = evalClassModel(svm_c, y_test1, y_pred_class)

    #Data for final graph
    #methodDict['Log. Regres.'] = accuracy_score * 100
```

In [48]:

SVM_C()

Support Vector Machine after feature selection -
Accuracy: 0.8058510638297872



Error: 0.19414893617021278
Precision: 0.7510204081632653
AUC Score: 0.7999433106575964
Cross-validated AUC: 0.8495489612974575

In [49]:

```
from sklearn.model_selection import RandomizedSearchCV

def tuningRandomizedSearchCV(model, param_dist):
    rand = RandomizedSearchCV(model, param_dist, cv=10, scoring='accuracy', n_iter=10, random_state=42)
    rand.fit(X1, y1)
    #rand.grid_scores_
```

In [50]:

```
def randomForest():
    # Calculating the best parameters
    forest = RandomForestClassifier(n_estimators = 100)

    featuresSize = feature_cols.__len__()
    param_dist = {"max_depth": [3, None],
                  "max_features": randint(1, featuresSize),
                  "min_samples_split": randint(2, 9),
```

```
# "min_samples_leaf": randint(1, 9),
  "criterion": ["gini", "entropy"]}
#tuningRandomizedSearchCV(forest, param_dist)


forest = RandomForestClassifier(max_depth = None, min_samples_leaf=8, min_samples_split=10)
my_forest = forest.fit(X_train1, y_train1)

y_pred_class = my_forest.predict(X_test1)

print('Random Forest after feature selection-')
accuracy_score = evalClassModel(my_forest, y_test1, y_pred_class, True)
```

In [51]: randomForest()

Random Forest after feature selection-
Accuracy: 0.7819148936170213



Error: 0.21808510638297873
Precision: 0.7544642857142857
AUC Score: 0.778344671201814
Cross-validated AUC: 0.8903591567565241

In []: